

# Containers

## Linux Kernel and Docker

Roberto Masocco  
`roberto.masocco@uniroma2.it`

University of Rome "Tor Vergata"  
Department of Civil Engineering and Computer Science Engineering  
Intelligent Systems Lab

May 5, 2022



**TOR VERGATA**  
UNIVERSITY OF ROME

School of Engineering

What follows is heavily based on specific features of the Linux kernel.  
**Compatibility with different platforms cannot be guaranteed.**

# Roadmap

1 Containers

2 Docker

3 Docker Compose

# Roadmap

1 Containers

2 Docker

3 Docker Compose

# Why Containers?

## Example: Packaging Applications

Suppose you are ready to distribute your new application:

- you need to be sure that it is compatible with all the **platforms** you chose to support;
- you need to figure out a way to deal with **dependencies**;
- you want to publish some kind of self-contained, easily-identifiable **package**.

# Why Containers?

## Example: Isolating Applications

Suppose you are deploying applications on a server:

- you want to define **resource quotas** and **permissions** for each;
- you want to be sure that each module has what it needs to operate, but **nothing more**;
- you want to **isolate** each module for security reasons, in case something goes wrong.

# Why Containers?

## Example: Replicating Environments

Suppose you are developing applications for a specific system (maybe with a different architecture):

- you want to have a **local copy** of such system without carrying one with you;
- you want to have all **libraries** and **dependencies** installed without tainting your own system;
- you would like to **deploy** the entire installation with just a few commands, without running any script but simply copying data.

# Why Containers?

A possible solution to many of the previous situations could be a set of **virtual machines**.

However, virtual machines are **slow**, hypervisors take up **system resources** and guest kernels must always be **tweaked**.

In each of the above scenarios something simpler would be enough, especially since **the OS is not involved**, only applications are.

This is what a **container** is.



Figure 1: FreeBSD jail logo



# Containers in the Linux kernel

Support for containers was added to the Linux kernel with a set of **features** starting from kernel 2.6 (2003), mainly:

- **control groups** (cgroups): defining different resource usage policies for groups of processes;
- **namespaces**: isolating processes and users in different "realms", both hardware (e.g. network stack) and software (e.g. PIDs);
- **capabilities**: defining what a process can do, with both hardware and software resources.

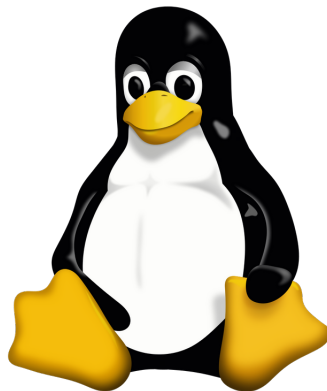


Figure 2: Tux

# Containers in the Linux kernel

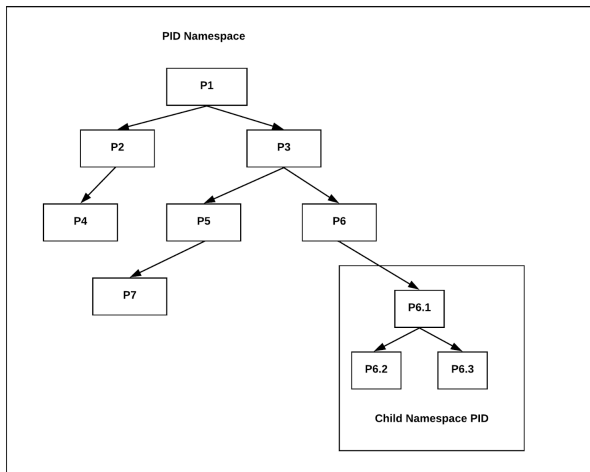


Figure 3: Nested PID namespaces

# Roadmap

1 Containers

2 Docker

3 Docker Compose

# Roadmap

1 Containers

2 Docker

3 Docker Compose